



MySQL and RabbitMQ

Or: OpenStack's core support softwares

Tyanko Aleksiev <tyanko.aleksiev@s3it.uzh.ch>

MySQL's role inside OpenStack

MySQL database assumes a core role in the set of support softwares inside OpenStack and is mainly used:

- for storing the state of the OpenStack cluster (VMs, users, volumes, tokens, etc),
- in the workflow of handling the various service requests (e.g. starting a new instance, creating a volume).

An alternative of MySQL in OpenStack could be any database supported by sqlalchemy.

RabbitMQ's role inside OpenStack

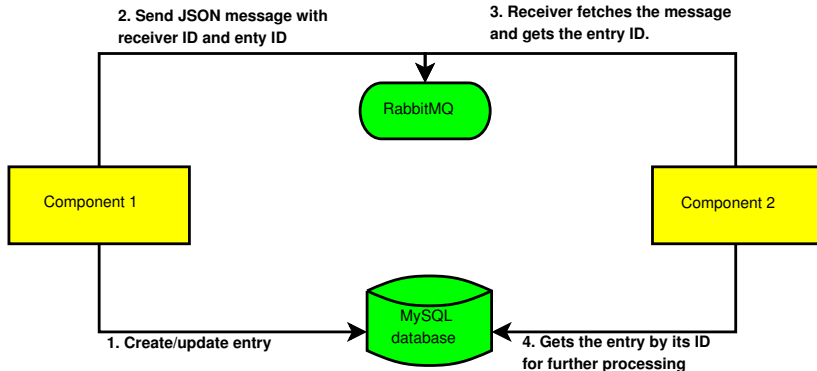
The *RabbitMQ* messaging system, used like Remote Procedure Call (RPC), is the second core support software and is mainly used:

- to provide a communication layer between the different stack components,
- as a consequence to enable the collaboration between the components.

The usage of RabbitMQ prevents the processes from polling the DB at every interaction.

Alternatives of RabbitMQ inside OpenStack are Qpid and ZeroMQ.

Example of RabbitMQ and MySQL usage during instance creation



Example of RabbitMQ and MySQL usage during instance creation

- a client sends a POST request to nova-api for a VM provisioning,
- nova-api writes an entry in the **DB**,
- nova-api posts then a JSON message in **RabbitMQ** queue (the message is for nova-scheduler),
- nova-scheduler reads the JSON message from **RabbitMQ**,
- it examines the overall cluster situation from the **DB**,
- nova-scheduler then decides which compute should start the VM and posts JSON message in its **RabbitMQ** queue,
- nova-compute gathers then the request information (from the **RabbitMQ** and **DB**) and proceeds with the VM provisioning.

MySQL High Availabiliti

MySQL is one of the most critical components because it contains the state of the whole Stack.

Having MySQL in HA becomes almost a requirement. OpenStack's recommended way to do it is by using the master - slave replication implementing DRBD and Pacemaker.

An active - active (N+1) alternative is using a Galera cluster which should become the state-of-the-art in the near future.

[This](#) link provides more information.

RabbitMQ High Availabiltiy

RabbitMQ can be configured in HA to prevent communication failure between the OpenStack components.

This can be achieved again, like in the case with MySQL, by using DRBD and Pacemaker.

An alternative method, by using again active - active (even if implementation is actually active - salve) mirrored servers.

[This](#) link provides detailed information.

Notes and Remarks

- Installation of MySQL and RabbitMQ is really straight forward and requires limited intervention.
- Configuration of RabbitMQ is limited to changing the default passwd and the bind address.
 - Conf. file is `/etc/rabbitmq/rabbitmq-env.conf`
 - Logs directory is: `/var/log/rabbitmq`
- Configuration of MySQL requires a little bit more than RabbitMQ.
 - Conf. file is: `/etc/mysql/my.cnf`
 - Logs directory is: `/var/log/mysql`
- We will see everything in more detail during the tutorial.